42

THE GAME OF 42 · FOR THE VISION FISH

The Hitchhiker’s Timeline

The strategy of the Game of 42: a proposed two-year timeline of events, convergences, and the focuses of the mice

DON’T PANIC

A proposal for the Vision Fish, set out by the mice within the Game of 42, a constructive instance of the Privacy Is Value model.

In brief

This is a proposal, not a paper. It sets out the mice's strategy for the next two years: the events we mean to hold, the convergences we are steering toward, and the focuses that occupy us in between. The prevailing trajectory of artificial intelligence points toward a single destination: one dominant model, one governing centre, and beneath both, one enormous lake into which the data of every person is gathered and mined. We are building the opposite: an architecture in which value is produced not by pooling data but by keeping boundaries. What follows names the central wager (a *poly-singularity*, many local sovereignties federating into one mutual without surrendering what makes them local), the dual-agent separation that makes such sovereignty workable, the trust graph that renders relationships verifiable, and the mutual-credit economy that prices contribution without extracting it. It then lays out the phased two-year plan, from a quiet declaration to a planetary deliberative structure, with the thresholds we are aiming at and an honest account of what is still unresolved. Read it as a plan we intend to act on.

How to read this

This is a plan, drawn from a run of working sessions, so the dates in it are intentions rather than commitments, and the headline figures (forty-two thousand users, $4.2M, forty-two billion tokens) are motifs and targets rather than forecasts. We hold one discipline throughout: where something is proven we say so, and where it is only hoped for we mark it as such. The formal treatment of the value model lives in a separate paper and is referenced, not reproduced, here. That model is Mitch's; the strategy that carries it is the work of the mice together.

1. The default future and its costs

The economics of contemporary machine intelligence reward concentration. Data accrues to whoever can pool it; capability accrues to whoever can pool the most; and the resulting advantage compounds. The endpoint that many participants now describe with equanimity is a centralised intelligence governing widening domains of life, fed by a commons of personal data that the people who generated it neither control nor benefit from.

The costs of this arrangement are not only the familiar ones of surveillance and lost privacy. They are structural. A single centre is a single point of capture, a single locus of failure, and a single account of what the future should be, imposed on a plural world that did not consent to it. Our wager is that the concentration is not inevitable, and that a plural alternative is both more just and, on the right accounting, more valuable. The rest of this document is the two-year path by which we mean to earn that claim.

2. Privacy as value

The core claim is an inversion. In the lake model, value is realised by pooling and mining data, and privacy appears as a cost to be managed down. In the model proposed here, value is realised by the opposite move: by keeping the boundary intact. Privacy ceases to be a friction on value creation and becomes its source.

This is the proposition Mitch has formalised as the Privacy Is Value model (developed across versions one through six, most recently the moving-ceiling R(t) line of work), which expresses a sovereignty value as a function of privacy maintained over time, with a small number of proven bounds and a larger number of clearly-labelled conjectures. The proven core is modest and deliberately so: a reconstruction ceiling, now treated as a moving bound R(t) rather than a fixed one, that holds under stated assumptions. The more ambitious properties, including optimality results, are marked as conjecture and are the subject of ongoing work. The discipline matters more than any single result: the aim is a value model whose claims a reviewer can separate cleanly into the demonstrated and the hoped-for.

From this inversion follows the destination. If boundaries are the source of value, then the right structure is not one boundary around everyone but many boundaries, each held locally, federating into a shared whole. This is the *poly-singularity*: many singularities, one commons, no centre. Everything that follows is the path that earns it.

3. The architecture of separation

Sovereignty at the level of an individual or a community is only tractable if the agent acting on your behalf is itself divided. The architecture rests on a separation between two roles: a boundary agent that protects and refuses (named in the project's mythology the Sword), and a connection agent that projects and relates (the Mage). Neither alone is safe. The boundary without the bond is a wall; the bond without the boundary is a flood. Held in tension, they make a sovereignty that can both keep its ground and reach beyond it. This separation is the grounding principle of the design, named the Separation Theorem; consistent with the hygiene stated above, its formal validation is in progress, and for now it should be read as a stated principle rather than a proven result.

A second separation runs between the local and the centralised. The architecture keeps an honest standoff between an agent that runs on your own machine, answering only to you, and the vast models in the cloud that can answer the largest questions but cannot be allowed to own you. The design does not resolve this by choosing a side. It builds for both and keeps the tension visible, because a system that pretends the tension away will quietly collapse toward the centre.

Trust, in this architecture, is not asserted but drawn. A *trust graph* overlays a record of what a party knows and has made (a knowledge graph) on a record of what they have staked and delivered (a promise graph). The relationship between the two layers is the trust, made computable. The pipeline runs from intent to promise to verified trust, and the middle term is never skipped. Promise Theory supplies the formal grammar: trust is built from voluntary, bilateral promises, not imposed obligations.

Two further structural claims are made, and both are marked as conjecture pending formal treatment: that the boundary of such a system can encode the volume it contains, after the manner of a holographic bound, so that a well-designed surface is sufficient to represent the bulk; and that the geometry of the sovereignty lattice admits substantial reductions in the cost of zero-knowledge verification. These are research directions, not settled results.

4. The economy

A plural architecture needs an economy that prices contribution without extracting it. The proposed instrument is a mutual-credit time bank with a layer most such systems lack. A participant pledges forty-two hours and receives forty-two credits; the ledger sums to zero; and threshold smart contracts refund automatically when a guide's target is not met, which removes the cold-start risk that ordinarily kills mutual-credit schemes. The hours are an honesty box: a senior professional and a child are credited at the same rate, which is a deliberate statement about what the system values.

Growth and governance share one geometry, the Game of 42. A participant begins alone, finds complementary roles, reaches a minimal working group of six, and then buds: seven groups of six compose a board of forty-two, and growth beyond that buds a further forty-two rather than inflating the first. Diversity of age, belief, and temperament is not an aspiration here but a load-bearing requirement, since an assembly that lacks it does not function as intended.

Identity, in turn, is economic. Each guide becomes real by crossing a threshold on one of three bonding curves: money, hours, or trust. The three together constitute the guide's identity, and the construction carries a security property as a side effect, since a guide that is private on one axis must be public on another and so cannot be attacked on both at once. Around this sit the conventional rails, treated with deliberate restraint: a crypto-economic design, a patron mechanism, a fair token launch with roughly a year of runway, and a redistributive fund. The standing economic discipline is stated plainly: bad money is bad money, to be vetted and, where necessary, declined regardless of size.

5. Governance and plurality

Governance is the point at which plurality stops being rhetoric and becomes structure. The proposed form is a nested deliberative body, forty-two assemblies by forty-two regions by forty-two sub-regions, in which seats are extended not only to people but to the parties usually excluded: nature, the young, the indigenous, and the unrecognised nations. Participants claim their own biome from the bottom up. The method is deliberation rather than polling: a selected, trust-graphed assembly meets across a full year and produces considered answers, not aggregated opinions.

The same structure resolves the hardest political problem in the design, which is the relationship to the frontier laboratories. The architecture cannot credibly position itself as an alternative to those laboratories while simultaneously courting them as partners. The resolution is one of sequence and ceremony: establish sovereign capacity first, then route any partnership through the citizen assemblies and an explicit trust ceremony before it is allowed near the commons. Donated compute, if it comes, enters as a bequest to a neutral third space, not as a lever of control.

6. The two-year plan

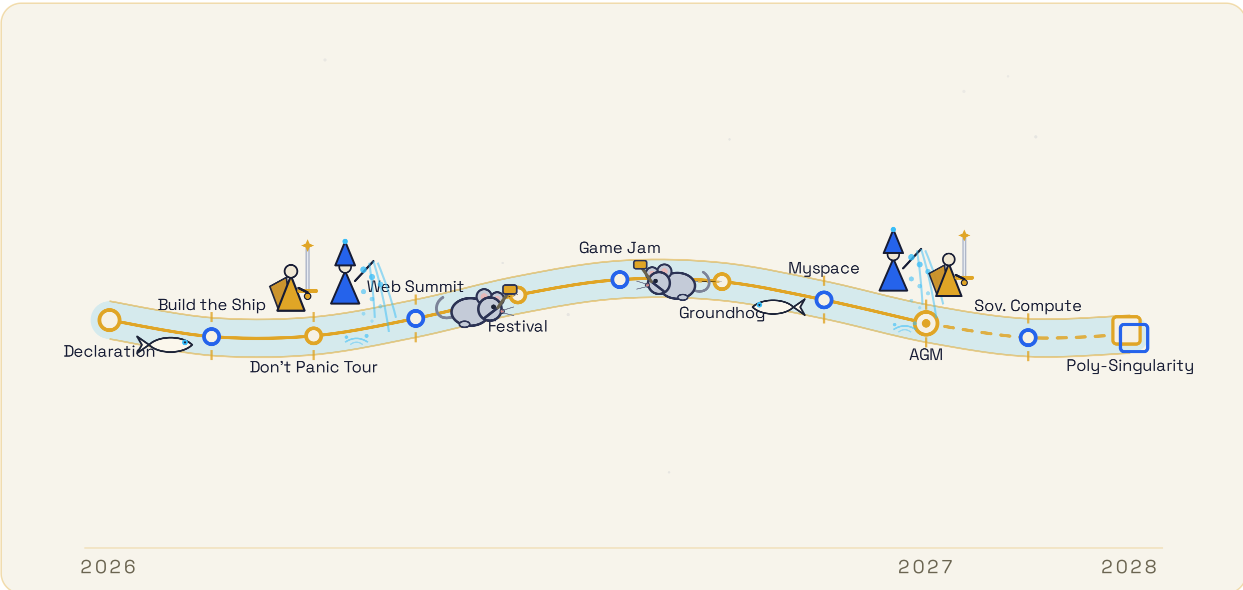

The two-year path, from the Declaration to the poly-singularity. The nodes are the milestones and the river is the road the work is laid along (solid where built, dashed across the 2028 horizon). The Swordsman sets the banks that bound the river and the Mage casts the water that fills them; the mice build the edges, and the Vision Fish swim the current.

The plan has three kinds of moving part, and it helps to name them: the *events* we hold (the milestones along the river), the *convergences* we steer toward (the AGM in year two, and the poly-singularity beyond it), and the *focuses* that occupy the mice between events (the systems we are building in the quieter months). What follows sets all three out in order.

The plan is not asserted but derived, and the method that derives it is worth stating because it is one of our more transferable assets. Every milestone is put through five framing questions (who cares; who is involved; what changes; what is the proof; and what makes it feel real), and only those that can answer the fifth survive. The plan is then back-cast from a fixed convergence point rather than extrapolated forward, so that each step is chosen for where it must lead rather than for where it happens to follow. And the whole is paced to grow at the speed of trust, on the principle that a movement which outruns its own trust is not a movement but an incident.

The milestones, in order, are as follows.

1. **Declaration (July 2026).** A quiet, unadvertised release of the founding text in a multi-perspectival form, measured by serendipity rather than reach, defended from the outset against impersonation and opportunist coinage.
2. **Building the ship (August 2026).** The unglamorous infrastructure (a mutual-credit time bank and its accounting) and its public counterpart (a live trust-graph visualisation), the first point at which the thesis becomes visible.
3. **The autumn season (September to December 2026).** A tour built on local events, student energy, and continuous broadcast; a first appearance at a major conference with on-chain financial privacy demonstrated as a positive-use case; and a distributed festival whose product is a tone, reframing a calm refusal to panic as a governance posture.
4. **Winter and the game jam (November 2026 to January 2027).** A fair token launch, a run of builder hackathons, and an open challenge to construct on a shared substrate.
5. **Institutional turn (February 2027).** A deliberate pivot to institutional partners and their risk managers, at which point commitments become capital and the project reads as a serious enterprise.
6. **The cultural turn (spring 2027).** The interface milestone at which participants compose their own front ends onto one substrate, alongside events that establish the annual cycle and make room for those who reject generated work.
7. **The convergence (July 2027).** A single distributed assembly across forty-two venues, at which the founding work passes into mutual ownership.
8. **The horizon (late 2027 into 2028).** Sovereign compute as a federated alternative, and then the two-faced arrival of the poly-singularity: a governance face (a respected planetary deliberative structure) and an answers face (an open, credentialed working of the largest questions).

Two success thresholds anchor the plan. By the convergence in July 2027: forty-two thousand participants and $4.2M committed to redistribution. By the cycle that follows in 2028: 4.2M participants and $42M. These figures are targets that shape the design, not predictions of outcome.

7. Open problems

The following problems are unresolved and are stated here rather than concealed, because they are where external judgement is most valuable.

- **Environmental cost.** The water and energy footprint of large-scale compute is the chief unresolved critique and does not yet have an adequate answer.
- **Frontier-laboratory terms.** Partnership is addressed by sequence and ceremony, but acquisition terms remain open.
- **Governance of donated resources.** The allocation of any donated compute must be settled before success arrives, not after.
- **Legality of incentives.** Lottery and prize mechanics are workable in some jurisdictions and a stretch in others.
- **Capital discipline.** Vigilance against extractive money and equity capture is a standing requirement, not a solved problem.
- **Provider strategy.** The choice between frontier partnership and a strict local-first posture is deferred and must eventually be made.
- **Representation.** The trust-graph visualisation must convey meaning rather than ornament, and its design is deliberately left open.

8. In sum

Our claim is narrow and, we think, important. If value can be made to flow from boundaries kept rather than data pooled, then the concentration that currently looks inevitable is a choice, and a worse one. What we are proposing instead is a federation of local sovereignties, each holding its own boundary, its own model, and its own voice, composing into a single mutual that can address the largest questions without anyone surrendering their ground. We have a word for the condition we are aiming at, borrowed and kept deliberately lowercase: eutierria, the state of belonging to a living system that has kept your boundary intact. This document is the two-year strategy by which we mean to build toward it, set out with what we are sure of and what we are not kept plainly apart. We are not asking you to admire the plan. We are asking you to test it, and to help us walk it.

Don't Panic.